

Annex IV Technical Documentation

Regulation: EU 2024/1689 Annex IV (Art. 11) **Variant:** full **Generated:** 2026-04-22T12:27:51.311199+00:00 **Status:** draft

System: external/retail-credit-risk-scorer · Retail Credit Risk Scorer · v3 · risk class high

1. General description

- **Intended purpose:** Evaluates consumer credit applications for a German retail bank. Inputs an application JSON (demographics, employment, income, existing obligations) and returns a PD score plus recommendation.
- **Provider:** Acme Bank GmbH
- **Contact:** ai-governance@acme.example
- **System identifier:** external/retail-credit-risk-scorer
- **Version:** v3
- **Form:** Managed runtime (API). Executed as an ephemeral server task on agents.renemurrell.de.
- **Underlying model:** anthropic/claude-sonnet-4-6
- **Required tools:** sql
- **User interface summary:** Operated via REST API (POST /v1/tasks) or CLI (agents-cli run). A browser dashboard surfaces status, results, and evidence packs.
- **IFU reference:** capability-card@v3
- **Third-country variants:** [PROVIDER: list Member-State-specific variants, if any; otherwise 'none']

2. Detailed description

- **Development methodology:** Shadow deployment on 3 months of historical applications before live cutover. Weekly ML-ops review. Change-control board approves every card-version increment.

Third-party components:

- Foundation model: anthropic/claude-sonnet-4-6
- Tool: sql
- Runtime: agents.renemurrell.de managed runtime
- **System architecture:** Each task runs on an isolated ephemeral server provisioned from a pre-baked snapshot, executed by the agent runtime, then destroyed. State is persisted to the platform event log, not to the compute node.

Data requirements:

- **input_schema:** [{ 'name': 'application', 'type': 'object', 'required': True}]
- **output_schema:** [{ 'name': 'score.json', 'type': 'json', 'guaranteed': True}]
- **constraints:** {}

- **training_data_summary:** Foundation model used without fine-tuning. Prompts parameterised with application fields only; no historical outcome data sent to the LLM.
- **Validation and testing:** [PROVIDER: test corpora, acceptance criteria, drift monitoring]
- **Pre-determined changes:** [PROVIDER: changes that do not trigger re-assessment, per Art. 43]
- **Continuous-learning behaviour:** No online learning. The agent is stateless between executions; model weights are fixed per capability-card version.

3. Monitoring information

- **Capabilities:** Evaluates consumer credit applications for a German retail bank. Inputs an application JSON (demographics, employment, income, existing obligations) and returns a PD score plus recommendation.

Known limitations:

- Degraded performance on applications with missing income data
- Not validated for non-retail (SME) lending

Observed accuracy metrics:

- **total_executions:** 0
- **success_count:** 0
- **observed_success_rate:** None
- **trust_score:** None
- **avg_duration_seconds:** None

Foreseeable unintended outcomes:

- [PROVIDER: incorrect outputs due to ambiguous inputs]
- [PROVIDER: privacy leakage if sensitive data is included in the prompt]
- [PROVIDER: bias reinforcement if the underlying model is biased]

Human oversight measures:

- Approval gates for high-risk agents (Art. 14)
- Immutable event log per execution (Art. 12)
- Signed capability cards with declared capabilities (Art. 13)
- Art. 50 provenance header on every output

4. Appropriateness of performance metrics

Platform-provided metrics:

- 5-factor trust score (success rate, duration accuracy, volume, recency, ratings)
- Observed success rate and failure-class breakdown from the event log
- Execution latency (avg and p95) per agent version

- **Metric rationale:** Approval/denial accuracy vs. a held-out human-reviewed benchmark, alongside demographic parity checks, are the most load-bearing metrics for our fairness-first framing.
- **Robustness testing:** [PROVIDER: adversarial inputs, edge cases, redundancy behaviour]

Cybersecurity testing:

- 3-path secret isolation (consumer / provider / platform)
- Ephemeral compute: no persistent tenant state
- Firewalled VPC, no SMTP egress

5. Risk management system

- **Platform risk class:** high
- **RMS owner:** Head of AI Governance

Identified risks:

- Disparate impact
- Prompt injection via free-text application notes
- Model drift after retraining of underlying LLM

Mitigation measures:

- Input sanitisation before prompt construction
- Fortnightly fairness dashboard review
- Immediate rollback if trust score drops >0.15
- **Review cadence:** monthly

6. Changes

- **Current version:** v3
- **Version history policy:** Each published change increments the agent capability card version. The platform stores prior versions and keeps their event logs for the Art. 19 retention window.
- **Substantial modification definition:** [PROVIDER: define what constitutes a substantial modification per Art. 43]

7. Applied standards

Harmonised standards applied:

- ISO/IEC 42001 as baseline
- EN ISO 9001 for engineering quality

Common specifications applied:

- (none)

Other adopted solutions:

- ISO/IEC 42001 (AI management system) as baseline reference where applicable

8. Declaration of conformity

- **DoC reference:** [PROVIDER: attach EU Declaration of Conformity for external/retail-credit-risk-scorer v3 per Annex V]
- **CE marking affixed:** False
- **EU database registration (Art. 49):** [PROVIDER: EU database registration identifier, per Art. 49]

9. Post-market monitoring

- **Plan summary:** [PROVIDER: post-market monitoring plan per Art. 72: data collection, analysis, re-assessment triggers]
- **Incident reporting process (Art. 73):** Serious incidents reported to the MSA within Art. 73 deadlines (15 days standard, 10 days critical infrastructure, 2 days widespread).

Feedback channels:

- Consumer ratings and complaint forms on the platform
 - Trust-score drift monitoring
 - [PROVIDER: additional channels: customer support, user surveys, etc.]
-

Generated by the agents.renemurrell.de Annex IV scaffolding engine. Sections marked [PROVIDER: ...] require completion by the placing-on-market organisation. Retain this document for 10 years per Art. 18.